

-- Create a database

```
CREATE DATABASE IF NOT EXISTS student_db;
```

-- Use the database

```
USE student_db;
```

-- Create a table for storing student information

```
CREATE TABLE IF NOT EXISTS student (
```

```
    student_id INT,
```

```
    name STRING,
```

```
    age INT,
```

```
    major STRING
```

```
)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY ','
```

```
STORED AS TEXTFILE;
```

-- Insert data into the table

```
INSERT INTO TABLE student VALUES
```

```
    (1, 'Alice', 20, 'Computer Science'),
```

```
    (2, 'Bob', 22, 'Mathematics'),
```

```
    (3, 'Charlie', 21, 'Physics');
```

-- Alternatively, use a LOAD command to add data from an external file

```
LOAD DATA LOCAL INPATH '/path/to/student_data.csv' INTO TABLE student;
```

```
SELECT * FROM student;
```

-- Find students majoring in Computer Science

```
SELECT * FROM student WHERE major = 'Computer Science';
```

-- Find students older than 21

```
SELECT * FROM student WHERE age > 21;
```

Update (U): Modifying data in the table

Hive doesn't support direct updates to tables in traditional ways (like SQL's `UPDATE`), but you can achieve it using an **INSERT OVERWRITE** approach:

-- Update the major of a student (e.g., Bob to 'Statistics')

```
INSERT OVERWRITE TABLE student
```

```
SELECT
```

```
    student_id,
```

```
    name,
```

```
    age,
```

```
    CASE
```

```
        WHEN name = 'Bob' THEN 'Statistics'
```

```
        ELSE major
```

```
    END AS major
```

```
FROM student;
```

Delete (D): Deleting data from the table

Hive doesn't support `DELETE` in the traditional sense. Instead, you can filter out the data you want to keep and overwrite the table:

-- Delete a student (e.g., remove Charlie)

```
INSERT OVERWRITE TABLE student
```

```
SELECT * FROM student WHERE name != 'Charlie';
```

```
DROP TABLE IF EXISTS student;
```

Create (C): Using Collection Types in Table Creation

```
CREATE TABLE student_with_collections (  
    student_id INT,  
    name STRING,  
    age INT,  
    major STRING,  
    grades ARRAY<INT>,          -- An array of grades  
    contact_info MAP<STRING, STRING>, -- A map for contact details (e.g., phone and email)  
    address STRUCT<city:STRING, zip:INT> -- A struct for address details  
)  
  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
COLLECTION ITEMS TERMINATED BY '|'  
MAP KEYS TERMINATED BY ':'  
STORED AS TEXTFILE;
```

-- Insert data into the table

```
INSERT INTO TABLE student_with_collections VALUES  
  
    (1, 'Alice', 20, 'Computer Science', ARRAY(85, 90, 95), MAP('phone', '1234567890', 'email',  
    'alice@example.com'), NAMED_STRUCT('city', 'New York', 'zip', 10001)),  
  
    (2, 'Bob', 22, 'Mathematics', ARRAY(78, 88, 92), MAP('phone', '9876543210', 'email',  
    'bob@example.com'), NAMED_STRUCT('city', 'Los Angeles', 'zip', 90001)),  
  
    (3, 'Charlie', 21, 'Physics', ARRAY(80, 85, 90), MAP('phone', '5551234567', 'email',  
    'charlie@example.com'), NAMED_STRUCT('city', 'Chicago', 'zip', 60601));
```

Querying Data from Collection Types

```
SELECT * FROM student_with_collections;
```

-- Retrieve the first grade of each student

```
SELECT name, grades[0] AS first_grade FROM student_with_collections;
```

-- Filter students with a grade above 90

```
SELECT * FROM student_with_collections WHERE ARRAY_CONTAINS(grades, 90);
```

-- Retrieve the city of each student

```
SELECT name, address.city AS city FROM student_with_collections;
```

-- Filter students based on ZIP code

```
SELECT * FROM student_with_collections WHERE address.zip = 90001;
```

```
INSERT OVERWRITE TABLE student_with_collections
```

```
SELECT
```

```
    student_id,
```

```
    name,
```

```
    age,
```

```
    major,
```

```
    CASE
```

```
        WHEN name = 'Alice' THEN ARRAY(88, 92, 96)
```

```
        ELSE grades
```

```
    END AS grades,
```

```
    contact_info,
```

```
    address
```

```
FROM student_with_collections;
```

```
INSERT OVERWRITE TABLE student_with_collections
```

```
SELECT * FROM student_with_collections WHERE name != 'Charlie';
```

```
-- Explode grades into multiple rows for analysis
```

```
SELECT name, grade
```

```
FROM student_with_collections LATERAL VIEW EXPLODE(grades) grade_table AS grade;
```

```
-- Retrieve all keys and values in the contact_info map
```

```
SELECT name, contact_key, contact_value
```

```
FROM student_with_collections
```

```
LATERAL VIEW EXPLODE(contact_info) map_table AS contact_key, contact_value;
```

```
-- Filter students based on city name
```

```
SELECT * FROM student_with_collections WHERE address.city = 'New York';
```

Partitioning in Hive

Partitioning organizes data into separate directories based on a column (or multiple columns). This helps improve query performance by scanning only the relevant partitions.

Creating a Partitioned Table

We'll partition the `student` table by the `major` column

```
CREATE TABLE student_partitioned (
```

```
    student_id INT,
```

```
    name STRING,
```

```
    age INT
```

```
)
```

```
PARTITIONED BY (major STRING)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY ','
```

```
STORED AS TEXTFILE;
```

Inserting Data into the Partitioned Table

You must specify the partition value when inserting data.

```
-- Insert data into the "Computer Science" partition
```

```
INSERT INTO TABLE student_partitioned PARTITION (major='Computer Science')
```

```
VALUES (1, 'Alice', 20);
```

```
-- Insert data into the "Mathematics" partition
```

```
INSERT INTO TABLE student_partitioned PARTITION (major='Mathematics')
```

```
VALUES (2, 'Bob', 22);
```

```
-- Insert data into the "Physics" partition
```

```
INSERT INTO TABLE student_partitioned PARTITION (major='Physics')
```

```
VALUES (3, 'Charlie', 21);
```

```
-- Show all partitions in the table
```

```
SHOW PARTITIONS student_partitioned;
```

```
-- Fetch all students from the "Computer Science" partition  
SELECT * FROM student_partitioned WHERE major = 'Computer Science';
```

```
LOAD DATA LOCAL INPATH '/path/to/computer_science_students.csv'  
INTO TABLE student_partitioned PARTITION (major='Computer Science');
```

Bucketing in Hive

Bucketing divides data into smaller, fixed-size files based on a column's hash function, often improving query performance for operations like joins.

Creating a Bucketed Table

We'll bucket the `student` table by the `student_id` column into 4 buckets.

```
CREATE TABLE student_bucketed (  
    student_id INT,  
    name STRING,  
    age INT,  
    major STRING  
)  
CLUSTERED BY (student_id) INTO 4 BUCKETS  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;
```

```
SET hive.enforce.bucketing = true;
```

```
-- Insert data into the bucketed table
```

```
INSERT INTO TABLE student_bucketed  
VALUES  
    (1, 'Alice', 20, 'Computer Science'),
```

```
(2, 'Bob', 22, 'Mathematics'),  
(3, 'Charlie', 21, 'Physics'),  
(4, 'David', 23, 'Statistics'),  
(5, 'Eve', 19, 'Biology');
```

```
-- Fetch students based on student_id  
SELECT * FROM student_bucketed WHERE student_id = 1;
```

You can combine **partitioning** and **bucketing** for better data organization and query optimization.

```
CREATE TABLE student_partitioned_bucketed (  
    student_id INT,  
    name STRING,  
    age INT  
)  
PARTITIONED BY (major STRING)  
CLUSTERED BY (student_id) INTO 3 BUCKETS  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;
```


-- Insert data into the "Computer Science" partition and bucket by student_id

```
INSERT INTO TABLE student_partitioned_bucketed PARTITION (major='Computer Science')
```

```
VALUES (1, 'Alice', 20), (4, 'David', 23);
```

-- Fetch students from a specific partition

```
SELECT * FROM student_partitioned_bucketed WHERE major = 'Computer Science';
```

-- Fetch a specific student within a partition

```
SELECT * FROM student_partitioned_bucketed WHERE major = 'Computer Science' AND student_id = 1;
```

-- Show partitions in the table

```
SHOW PARTITIONS student_partitioned_bucketed;
```

-- To check the files in a specific partition (use shell command on the Hadoop file system)

```
hadoop fs -ls /path/to/hive/warehouse/student_partitioned_bucketed/major=Computer Science/
```

🎬 **Partitioning** improves query performance by dividing data into directories based on column values.

🎬 **Bucketing** splits data into smaller files within partitions, enabling faster joins and better data distribution.